

VISUAL SERVOING FOUNDATION-MODEL HYBRID: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study visual servoing and foundation policies. The motivating bet is: Define when classical servo loops should override learned action decoding.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For visual servoing and foundation policies, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2022; 2025a; 2009; 2025b; 2019; 2024). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.588	0.412
observed_only_wam	hidden_mode_shift	0.163	0.838
augmentation_as_data_wam	nominal	0.625	0.375
augmentation_as_data_wam	hidden_mode_shift	0.292	0.708
scalar_uncertainty_wam	nominal	0.703	0.297
scalar_uncertainty_wam	hidden_mode_shift	0.333	0.667
proposed_branch_mechanism	nominal	0.787	0.213
proposed_branch_mechanism	hidden_mode_shift	0.568	0.432

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Visual Servoing Foundation-Model Hybrid argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- A hybrid visual servo controller for robust grasping by wheeled mobile robots, 2009. <https://doi.org/10.1109/tmech.2009.2034740>.
- Picking robot visual servo control based on modified fuzzy neural network sliding mode algorithms, 2019. <https://doi.org/10.3390/electronics8060605>.
- Constrained image-based visual servoing of robot manipulator with third-order sliding-mode observer, 2022. <https://doi.org/10.3390/machines10060465>.
- Adaptive visual servoing for the robot manipulator with extreme learning machine and reinforcement learning, 2024. <https://doi.org/10.1002/asjc.3208>.
- Hybrid deep learning framework for eye-in-hand visual control systems, 2025a. <https://doi.org/10.3390/robotics14050066>.
- Hybrid quantum-classical algorithms for high-precision sensor fusion in robotics applications, 2025b. <https://doi.org/10.62802/7sgy8455>.